

# MES/SCADA RAG System - Vibe Coding Guide 🚀

*The ultimate development philosophy and standards for building production-ready industrial RAG systems*

## 🎯 THE VIBE

**We build systems that work in the real world, not demos that impress in meetings.**

- **Architecture First:** No code without clear design. Think twice, code once.
- **Domain-Driven:** Start from business logic, not cool technology
- **Contract-First:** API endpoints defined before implementation
- **Production-Ready:** Every line of code should be ready for 24/7 industrial environments
- **Documentation-Driven:** Future developers (including yourself) will thank you

## 🏗️ SYSTEM OVERVIEW

🏭 Manufacturing Hierarchy (ISA-95 inspired, customer flexible)

↓

📄 Document Management (PDF, DOCX, XLSX processing)

↓

🧠 RAG Engine (Local embeddings + Vector search)

↓

💬 Chat Interface (Conversational AI for operators)

**Stack:** FastAPI + PostgreSQL + Qdrant + Ollama + Vanilla JS + Docker Compose

## 🎨 CODE STYLE - NON-NEGOTIABLE

### Python Standards

python

# ✅ Good - Type hints mandatory

```
def process_document(  
    file_path: str,  
    metadata: Dict[str, Any],  
    node_ids: List[int]  
    ) -> Optional[Document]:  
    """Process document with proper error handling."""  
    pass
```

# ❌ Bad - No types, unclear purpose

```
def process(file, data, ids):  
    pass
```

## Naming Conventions

python

*# Variables & functions: snake\_case*

```
user_documents = get_user_documents()
```

*# Classes: PascalCase*

```
class DocumentProcessor:  
    pass
```

*# Constants: UPPER\_SNAKE\_CASE*

```
MAX_FILE_SIZE = 50 * 1024 * 1024 # 50MB
```

*# Private methods: leading underscore*

```
def _internal_helper():  
    pass
```

## Error Handling - Industrial Grade

python

*#  Custom exceptions with context*

```
class DocumentProcessingError(Exception):  
    def __init__(self, message: str, file_path: str, original_error: Exception = None):  
        self.file_path = file_path  
        self.original_error = original_error  
        super().__init__(message)
```

*#  Structured logging*

```
logger.error(  
    "Document processing failed",  
    document_id=document_id,  
    error=str(e),  
    operation="process_document"  
)
```

## PROJECT STRUCTURE - SACRED GEOMETRY

```

src/
├── api/                # FastAPI routes
│   └── v1/
│       ├── hierarchy.py # Manufacturing hierarchy endpoints
│       ├── documents.py # Document management
│       ├── search.py    # RAG search endpoints
│       └── auth.py      # Authentication
├── core/               # Business logic (domain layer)
│   ├── hierarchy/
│   ├── documents/
│   ├── search/
│   └── auth/
├── models/             # SQLAlchemy models
├── schemas/            # Pydantic models
├── services/           # External service integrations
├── utils/              # Utilities & helpers
└── config.py           # Configuration management

```

## DATABASE DESIGN PRINCIPLES

### Core Tables

```

sql

-- Flexible hierarchy (not strict ISA-95)
hierarchy_nodes:
- id, name, code, description
- parent_id (self-reference)
- level_type (varchar, not enum - customer flexibility)
- custom_attributes (JSON for extensions)

-- Document management
documents:
- id, filename, file_path, content_hash
- metadata (JSON), created_at, updated_at
- processing_status, content_extracted

-- Many-to-many relationships
document_node_associations:
- document_id, node_id, relationship_type

```

### Database Operations - Always Safe

python

#  *Dependency injection + proper error handling*

```
def get_document_by_id(db: Session, document_id: int) -> Optional[Document]:  
    try:  
        return db.query(Document).filter(Document.id == document_id).first()  
    except SQLAlchemyError as e:  
        logger.error(f"Database error fetching document {document_id}: {e}")  
        raise DatabaseError(f"Failed to fetch document: {e}")
```

#  *Context manager for sessions*

```
@contextmanager  
def get_db_session() -> Session:  
    db = SessionLocal()  
    try:  
        yield db  
        db.commit()  
    except Exception:  
        db.rollback()  
        raise  
    finally:  
        db.close()
```

## RAG PIPELINE - THE MAGIC

### Document Processing Flow

python

# 1. Text Extraction (format-aware)

PDF → PyMuPDF (fast, reliable)

DOCX → python-docx

XLSX → pandas + openpyxl

TXT → charset detection

# 2. Semantic Chunking

- Chunk size: 512-1024 tokens
- Overlap: 50-100 tokens
- Preserve document structure metadata

# 3. Embeddings Generation

- Local Ollama (privacy + cost control)
- Model: sentence-transformers/all-MiniLM-L6-v2
- Fallback to cloud embeddings if needed

# 4. Vector Storage

- Qdrant with metadata filtering
- Hierarchical context preservation

## Search Algorithm

python

```
def rag_search(query: str, user_hierarchy_context: List[int]) -> SearchResults:
    """
    1. Query preprocessing (cleaning, expansion)
    2. Generate query embeddings
    3. Vector similarity search in Qdrant
    4. Apply hierarchy filters based on user permissions
    5. Rank results by relevance + context
    6. Format response with citations
    """
    pass
```

## SECURITY - ENTERPRISE GRADE

### Authentication Strategy

python

*# Multi-provider support*

```
providers = {  
    'local': LocalAuthProvider(),    # Username/password fallback  
    'saml': SAMLProvider(),          # Enterprise SSO  
    'oidc': OIDCProvider(),          # Modern SSO  
    'ldap': ActiveDirectoryProvider() # Legacy AD  
}
```

*# JWT + Redis backing (revocation capability)*

JWT\_SECRET\_KEY = "your-secret-key"

REDIS\_SESSION\_PREFIX = "session:"

## Authorization Model

python

*# Role-based access control*

```
roles = {  
    'system_admin': ['*'], # Full access  
    'hierarchy_manager': ['hierarchy.*', 'documents.read'],  
    'document_manager': ['documents.*', 'hierarchy.read'],  
    'user': ['documents.read', 'search.*'] # Based on hierarchy assignment  
}
```



## DEPLOYMENT - ONE COMMAND MAGIC

### Docker Compose Services

yaml

services:

frontend: *# Nginx + static files*

api: *# FastAPI application*

postgres: *# Metadata database*

qdrant: *# Vector database*

redis: *# Cache & sessions*

ollama: *# Local LLM for embeddings*

nginx: *# Reverse proxy*

### File Storage Structure

```
/data/uploads/  
/{year}/    # 2025/  
/{month}/  # 01/  
/{hash}/   # abc123.../  
file.pdf    # original filename  
meta.json   # extracted metadata
```

## TESTING PHILOSOPHY

### Test Pyramid

```
python  
  
# Unit Tests - Business logic  
def test_hierarchy_node_creation():  
    node = create_hierarchy_node("Plant A", "plant", parent=None)  
    assert node.name == "Plant A"  
    assert node.level_type == "plant"  
  
# Integration Tests - API endpoints  
def test_document_upload_endpoint(client, auth_headers):  
    response = client.post("/api/v1/documents/",  
                           files={"file": test_pdf},  
                           headers=auth_headers)  
    assert response.status_code == 201  
  
# E2E Tests - Critical workflows  
def test_complete_rag_workflow():  
    # Upload document → Process → Search → Verify results  
    pass
```

## PERFORMANCE TARGETS

- **Response Time:** < 3s for RAG queries (95th percentile)
- **Throughput:** 100+ concurrent users
- **Document Processing:** < 30s per document
- **System Uptime:** > 99.5%
- **Search Accuracy:** > 90% relevant results

## COMMON PITFALLS TO AVOID

### Don't Do This

python

*# No error handling*

```
def process_file(file):  
    return extract_text(file)
```

*# Magic numbers everywhere*

```
if file_size > 52428800: # What is this number?  
    raise Exception("Too big")
```

*# Hardcoded paths*

```
with open("/var/data/file.txt"):  
    pass
```

## ✓ Do This Instead

python

*# Proper error handling with context*

```
def process_file(file_path: str) -> ProcessingResult:  
    try:  
        return extract_text(file_path)  
    except FileNotFoundError:  
        raise DocumentNotFoundError(f"File not found: {file_path}")  
    except PermissionError:  
        raise DocumentAccessError(f"Permission denied: {file_path}")
```

*# Named constants*

```
MAX_FILE_SIZE = 50 * 1024 * 1024 # 50MB  
if file_size > MAX_FILE_SIZE:  
    raise FileTooLargeError(f"File size {file_size} exceeds limit {MAX_FILE_SIZE}")
```

*# Configuration-driven paths*

```
file_path = config.UPLOAD_DIR / filename
```

## SPRINT SUCCESS CRITERIA

### Technical Milestones

- **Sprint 1:** Complete Docker Compose environment running
- **Sprint 3:** Basic CRUD operations working end-to-end
- **Sprint 5:** First successful RAG query with proper results
- **Sprint 7:** Complete user workflow functioning
- **Sprint 10:** Production-ready deployment with monitoring



## Quality Gates (ALL MUST PASS)

- ☐ All tests passing (unit + integration + e2e)
- ☐ Code coverage > 80%
- ☐ Security scan passing
- ☐ Performance benchmarks met
- ☐ Documentation complete
- ☐ Code review approved

## DEVELOPMENT WORKFLOW

### Git Flow

```
bash
```

```
# Feature branches
```

```
git checkout -b feature/sprint-X-feature-name
```

```
# Conventional commits
```

```
git commit -m "feat(documents): add PDF text extraction with error handling"
```

```
git commit -m "fix(search): resolve vector similarity calculation bug"
```

```
git commit -m "docs(api): update OpenAPI specification for auth endpoints"
```

### Code Review Checklist

- ☐ Type hints present and correct
- ☐ Error handling implemented
- ☐ Tests written and passing
- ☐ Logging added for debugging
- ☐ Documentation updated
- ☐ Security considerations addressed
- ☐ Performance impact assessed

## THE MANTRA

**"Make it work, make it right, make it fast, make it maintainable"**

1. **Work:** Solve the business problem first
2. **Right:** Clean, readable, well-architected code
3. **Fast:** Optimize for performance where it matters
4. **Maintainable:** Future developers will understand and extend it

---

*Remember: We're building systems for industrial environments where downtime costs money and operators need answers fast. Every line of code should reflect this reality.*

Now go build something awesome! 🛠️