

Technical Standards & Decision Log

MES/SCADA RAG System

ARCHITECTURAL DECISIONS RECORD (ADR)

ADR-001: Database Strategy

Date: 2025-01-XX

Status: DECIDED

Context: Potřebujeme kombinovat relační data (hierarchie, metadata) s vektorovými embeddings






Decision:

- Primary DB: PostgreSQL 15+ (relační data, metadata, audit)
- Vector DB: Qdrant (embeddings, similarity search)
- Cache: Redis (sessions, query cache)

Rationale:

- PostgreSQL: Mature, excellent JSON support, ACID compliance
- Qdrant: Production-ready, excellent metadata filtering, scalable
- Redis: Fast, proven session management

Consequences:

-  Proven technology stack
-  Good performance for expected scale
-  Better metadata filtering than ChromaDB
-  Multiple databases = complexity
-  Qdrant requires separate service management

ADR-002: Vector DB Decision

Date: 2025-07-XX

Status:  DECIDED

Context: Potřebujeme spolehlivou a škálovatelnou vektorovou databázi pro uložení embeddingů (např. texty, poznámky, biosignály) a provádění podobnostních dotazů v rámci AI funkcionalit (např. vyhledávání, doporučování, RAG).

Decision: Zvolená databáze: Qdrant

- Self-hosted, open-source vektorová DB

- REST API, gRPC, a Python-native SDK
- Podpora filtrování podle metadat
- Možnost nasazení lokálně i v cloudu

Rationale:

- ● Snadné nasazení (Docker, binary)
- ● Podpora metadatových filtrů a scoringu
- ● Vhodné pro RAG, recommendation, personalizaci
- ● Aktivní vývoj a dokumentace
- ● Nutnost samostatného běhu služby (na rozdíl od ChromaDB)

Consequences:

- ✓ Robustní základ pro AI vyhledávání
- ✓ Možnost škálování podle potřeb
- ✓ Možnost snadné výměny embeddingů, retrain
- ✗ Potřeba správy dalšího serveru (Docker/service)

ADR-003: File Storage Strategy

Date: 2025-01-XX

Status: DECIDED

Context: Stovky GB dokumentů, různé formáty, backup requirements

Decision: Filesystem-based storage s organizovanou strukturou

```
/data/uploads/  
/{year}/    # 2025/  
/{month}/  # 01/  
/{hash}/   # abc123.../  
file.pdf   # original filename  
meta.json  # extracted metadata
```

Rationale:

- Jednoduchost implementace a debugingu
- Standard filesystem backup tools
- Žádné vendor lock-in
- Easy migration k object storage později

Alternatives Considered:

- MinIO: Overkill pro MVP, ale good migration path
- Database BLOB storage: Performance issues při velkých souborech

ADR-004: Manufacturing Hierarchy Model

Date: 2025-01-XX

Status: DECIDED

Context: ISA-95 je standard, ale zákazníci potřebují flexibility

Decision: Flexibilní hierarchický model inspirovaný ISA-95

sql

hierarchy_nodes:

- id, name, code, description
- parent_id (self-reference)
- level_type (**varchar**, ne **enum**)
- level_order (hierarchy position)
- custom_attributes (JSON)

Rationale:

- ISA-95 jako default template, ale ne enforcement
- Zákazníci mohou definovat vlastní level types
- Zachována hierarchická struktura
- Extensible přes JSON attributes

ADR-005: Authentication Architecture

Date: 2025-01-XX

Status: DECIDED

Context: Enterprise SSO requirements + local fallback

Decision: Multi-provider authentication s JWT tokens

Implementation:

- JWT tokens (stateless)
- Redis session backing (revocation capability)
- Provider abstraction layer:
 - Local (username/password)
 - SAML 2.0 (enterprise SSO)
 - OIDC (modern SSO)

- Active Directory (LDAP)

Rationale:

- Flexibility pro různé zákazníky
- Standard protocols
- Stateless token = scalability
- Session backing = security

ADR-006: RAG Implementation Strategy

Date: 2025-01-XX

Status: DECIDED

Context: Balance mezi accuracy, privacy, a cost

Decision: Local embeddings s Ollama + hierarchical context

Pipeline:

1. Document text extraction
2. Semantic chunking (paragraph-aware)
3. Local embeddings generation (Ollama)
4. Qdrant storage s metadata
5. Hybrid search (vector + keyword + hierarchy filter)

Model Selection: sentence-transformers/all-MiniLM-L6-v2 (default)

- Good multilingual support
- Reasonable size/performance
- Proven v RAG applications

CODING STANDARDS

Python Code Standards

Project Structure

```

src/
├── api/           # FastAPI routes
│   ├── v1/
│   │   ├── hierarchy.py # Manufacturing hierarchy endpoints
│   │   ├── documents.py # Document management
│   │   ├── search.py    # RAG search endpoints
│   │   └── auth.py      # Authentication
│   └── dependencies.py  # Shared dependencies
├── core/          # Business logic (domain layer)
│   ├── hierarchy/
│   ├── documents/
│   ├── search/
│   └── auth/
├── models/        # SQLAlchemy models
├── schemas/       # Pydantic models
├── services/      # External service integrations
├── utils/         # Utilities & helpers
└── config.py      # Configuration management

```

Naming Conventions

python

Variables & functions: snake_case

user_name = "john_doe"

def get_user_documents():

pass

Classes: PascalCase

class DocumentProcessor:

pass

Constants: UPPER_SNAKE_CASE

MAX_FILE_SIZE = 50 * 1024 * 1024 # 50MB

Private methods: leading underscore

def _internal_helper():

pass

Database tables: plural snake_case

class HierarchyNode(Base):

__tablename__ = "hierarchy_nodes"

Type Hints (MANDATORY)

python

```
from typing import List, Optional, Dict, Any, Union
from pydantic import BaseModel
```

```
def process_document(
    file_path: str,
    metadata: Dict[str, Any],
    node_ids: List[int]
) -> Optional[Document]:
    """
    Process uploaded document and create database record.
```

Args:

file_path: Path to uploaded file
metadata: Document metadata dictionary
node_ids: List of hierarchy node IDs to associate

Returns:

Created Document instance or None on failure

Raises:

DocumentProcessingError: If file processing fails
ValidationError: If metadata validation fails

```
    """
```

```
    pass
```

Error Handling Standards

python

Custom exceptions

```
class DocumentProcessingError(Exception):  
    """Raised when document processing fails"""  
    def __init__(self, message: str, file_path: str, original_error: Exception = None):  
        self.file_path = file_path  
        self.original_error = original_error  
        super().__init__(message)
```

HTTP exception handling

```
from fastapi import HTTPException, status  
  
@app.exception_handler(DocumentProcessingError)  
async def document_processing_handler(request, exc):  
    return HTTPException(  
        status_code=status.HTTP_422_UNPROCESSABLE_ENTITY,  
        detail={  
            "message": str(exc),  
            "file_path": exc.file_path,  
            "type": "document_processing_error"  
        }  
    )
```

Database Operations

python

Always use dependency injection

from sqlalchemy.orm import Session

from fastapi import Depends

def get_document_by_id(db: Session, document_id: int) -> Optional[Document]:

"""Get document by ID with error handling"""

try:

return db.query(Document).filter(Document.id == document_id).first()

except SQLAlchemyError as e:

logger.error(f"Database error fetching document {document_id}: {e}")

raise DatabaseError(f"Failed to fetch document: {e}")

Always close sessions properly

@contextmanager

def get_db_session() -> Session:

db = SessionLocal()

try:

yield db

db.commit()

except Exception:

db.rollback()

raise

finally:

db.close()

Logging Standards

python

```
import logging
import structlog
```

```
# Structured logging setup
```

```
logger = structlog.get_logger(__name__)
```

```
def process_document(document_id: int):
```

```
    logger.info(
        "Starting document processing",
        document_id=document_id,
        operation="process_document"
```

```
    )
```

```
    try:
```

```
        # processing logic
```

```
        logger.info(
            "Document processing completed",
            document_id=document_id,
            processing_time_ms=processing_time
        )
```

```
    except Exception as e:
```

```
        logger.error(
            "Document processing failed",
            document_id=document_id,
            error=str(e),
            operation="process_document"
```

```
        )
```

```
        raise
```